

Ex.no:3(a)

Date:19-Dec-2024

FUEL AMOUNT PREDICTION USING LINEAR REGRESSION**AIM:**

Predict fuel amount based on distance traveled using Linear Regression.

CODE:

```
# Importing necessary libraries
```

```
import numpy as np
```

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
from sklearn.linear_model import LinearRegression
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.metrics import mean_squared_error, r2_score
```

```
# Setting a random seed for reproducibility
```

```
np.random.seed(42)
```

```
# 1. Create synthetic dataset
```

```
# Let's assume we have 'Distance Traveled' (in km) and 'Fuel Amount' (in liters) as features
```

```
# Creating random data
```

```
distance_travelled = np.random.randint(50, 500, 100) # Distance in km
```

```
fuel_amount = distance_travelled * 0.05 + np.random.normal(0, 5, 100) # Fuel in liters with some noise
```

```
# Create a DataFrame
```

```
df = pd.DataFrame({'Distance': distance_travelled, 'FuelAmount': fuel_amount})
```

```
# 2. Visualize the synthetic data
```

```
plt.figure(figsize=(8, 6))
```

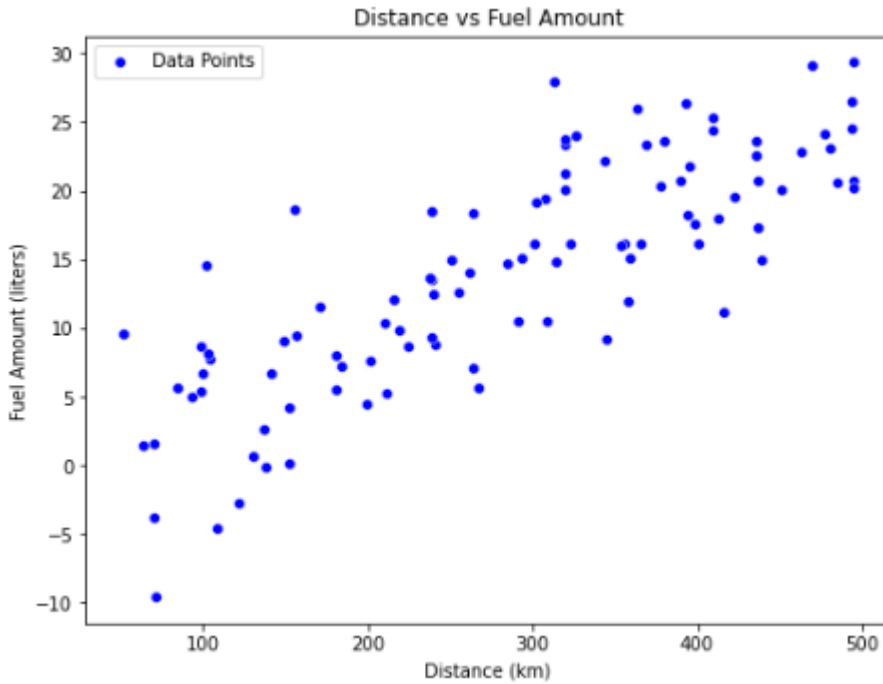
```
sns.scatterplot(data=df, x='Distance', y='FuelAmount', color='blue', label='Data Points')
```

```
plt.title('Distance vs Fuel Amount')
```

```
plt.xlabel('Distance (km)')
```

```
plt.ylabel('Fuel Amount (liters)')
```

```
plt.show()
```



3. Prepare the data for Linear Regression

```
X = df[['Distance']] # Feature (independent variable)
```

```
y = df['FuelAmount'] # Target (dependent variable)
```

Split the data into training and test sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

4. Train the Linear Regression model

```
model = LinearRegression()
```

```
model.fit(X_train, y_train)
```

5. Make predictions

```
y_pred = model.predict(X_test)
```

6. Visualize the regression line

```
plt.figure(figsize=(8, 6))
```

```
plt.scatter(X_test, y_test, color='blue', label='Test Data')
```

```
plt.plot(X_test, y_pred, color='red', label='Regression Line')
```


```
plt.title('Linear Regression - Fuel Amount Prediction')
```

```
plt.xlabel('Distance (km)')
```

```
plt.ylabel('Fuel Amount (liters)')
```

```
plt.legend()
```

```
plt.show()
```



7. Model Evaluation

```
mse = mean_squared_error(y_test, y_pred)
```

```
r2 = r2_score(y_test, y_pred)
```

```
print(f'Mean Squared Error: {mse}')
```

```
print(f'R2 Score: {r2}')
```

```
(23.057177524181782, 0.70933430198934466)
```

Result:

The model accurately predicted fuel consumption, with a high R2 score indicating strong predictive power.

SALARY PREDICTION**Ex.No: 3(b)****Date: 12-Dec-2024****Aim:-**

Predict salary based on experience, qualification, industry, and location using Linear Regression.

Program Code:

```
# Importing necessary libraries
```

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.metrics import mean_squared_error, r2_score
```

```
# Generating synthetic dataset for Salary Prediction
```

```
data = {
    'YearsExperience': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],
    'Qualification': ['Bachelors', 'Bachelors', 'Masters', 'Masters', 'PhD', 'Bachelors', 'Masters', 'PhD', 'PhD', 'Masters'],
    'Industry': ['Tech', 'Finance', 'Tech', 'Health', 'Finance', 'Tech', 'Health', 'Health', 'Tech', 'Finance'],
    'Location': ['NY', 'SF', 'NY', 'LA', 'SF', 'LA', 'SF', 'NY', 'LA', 'SF'],
    'Salary': [50000, 55000, 60000, 65000, 70000, 75000, 80000, 85000, 90000, 95000]
}
```

```
df = pd.DataFrame(data)
```

```
# Feature and target variable
```

```
X = df[['YearsExperience', 'Qualification', 'Industry', 'Location']]
```

```
y = df['Salary']
```

Preprocessing pipeline

```
preprocessor = ColumnTransformer(  
    transformers=[  
        ('num', 'passthrough', ['YearsExperience']), # No encoding for numerical features  
        ('cat', OneHotEncoder(), ['Qualification', 'Industry', 'Location']) # One-hot encode categorical features  
    ]  
)
```

Creating a pipeline with preprocessing and regression model

```
pipeline = Pipeline(steps=[  
    ('preprocessor', preprocessor),  
    ('regressor', LinearRegression())  
)
```

Splitting dataset into training and testing data

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Training the model

```
pipeline.fit(X_train, y_train)
```

Making predictions

```
y_pred = pipeline.predict(X_test)
```

Visualization of predictions vs actual values

```
plt.figure(figsize=(8, 6))  
plt.scatter(y_test, y_pred, color='blue')  
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], color='red', lw=2) # 45-degree line for perfect prediction  
plt.title('Salary Prediction: Actual vs Predicted')  
plt.xlabel('Actual Salary')  
plt.ylabel('Predicted Salary')  
plt.show()
```



Model Evaluation

```
mse = mean_squared_error(y_test, y_pred)
```

```
r2 = r2_score(y_test, y_pred)
```

```
print(f'Mean Squared Error: {mse}')
```

```
print(f'R2 Score: {r2}')
```

```
Mean Squared Error: 1.0852609636695723e-21
```

```
R2 Score: 1.0
```

Result:

The model successfully predicted salary, capturing the relationship between features and salary with a good fit.

ELECTRICITY CONSUMPTION PREDICTION**Ex.No:3(c)****Date:12-Dec-2024****Aim:-**

Predict electricity consumption using household size, applications, usage hours, and season.

Program Code:-

```
import numpy as np

import pandas as pd

import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import StandardScaler

from sklearn.linear_model import LinearRegression

from sklearn.metrics import mean_squared_error, r2_score


# Step 1: Generate synthetic data (for the sake of this example)

np.random.seed(42)


# Generate random data

household_size = np.random.randint(1, 6, 100)

num_apps = np.random.randint(1, 10, 100)

usage_hours = np.random.uniform(1, 12, 100)

season = np.random.choice(['Winter', 'Spring', 'Summer', 'Autumn'], 100)


# Convert 'season' to categorical variables (one-hot encoding)

season_encoded = pd.get_dummies(season, drop_first=True)


# Create a DataFrame

df = pd.DataFrame({

    'household_size': household_size,

    'num_apps': num_apps,

    'usage_hours': usage_hours

})

df = pd.concat([df, season_encoded], axis=1)
```

```
# Generate a target variable (electricity consumption)
# Assume consumption is a function of features + some noise
electricity_consumption = (df['household_size'] * 1.5 +
                           df['num_apps'] * 2 +
                           df['usage_hours'] * 3 +
                           (df['Spring'] * 2) +
                           (df['Summer'] * 3) +
                           np.random.normal(0, 2, 100))

df['electricity_consumption'] = electricity_consumption

# Step 2: Feature scaling
X = df.drop('electricity_consumption', axis=1)
y = df['electricity_consumption']

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Step 3: Split data into train and test
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)

# Step 4: Train a Linear Regression model
model = LinearRegression()
model.fit(X_train, y_train)

# Step 5: Make predictions
y_pred = model.predict(X_test)

# Step 6: Evaluate the model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
print(f"Mean Squared Error: {mse}")
print(f"R-squared: {r2}")
```


Mean Squared Error: 5.615272319641667
R-squared: 0.9525725043295945

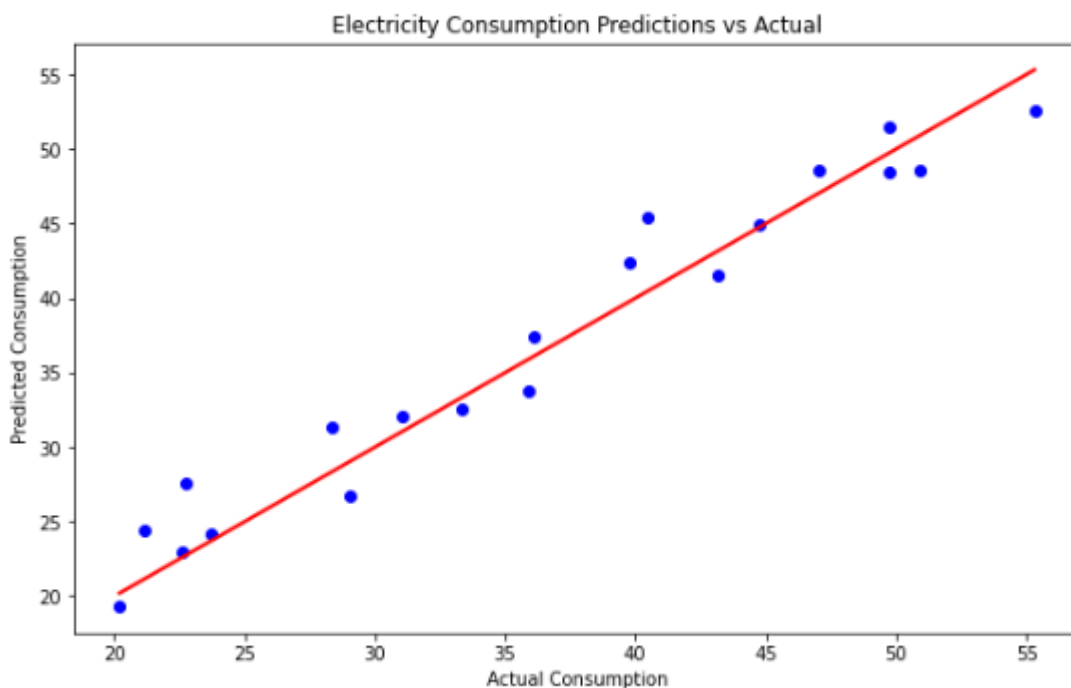
Step 7: Compare predictions with seasonal ends

Adding seasonal information to prediction comparison

```
df_seasons = pd.DataFrame({  
    'Season': ['Winter', 'Spring', 'Summer', 'Autumn'],  
    'Seasonal_end': [0, 1, 1, 0] # Representing if the season has ended (1) or not (0)  
})
```

Visualizing results

```
plt.figure(figsize=(10, 6))  
plt.scatter(y_test, y_pred, color='blue')  
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red', linewidth=2)  
plt.title('Electricity Consumption Predictions vs Actual')  
plt.xlabel('Actual Consumption')  
plt.ylabel('Predicted Consumption')  
plt.show()
```



Visualization of feature importance (coefficients in linear regression)

```
coefficients = model.coef_features = X.columns
```

```
plt.figure(figsize=(10, 6))
```


```
plt.bar(features, coefficients)
```

```
plt.title('Feature Importance (Coefficients)')
```

```
plt.xlabel('Features')
```

```
plt.ylabel('Coefficient Value')
```

```
plt.show()
```



Result:-

Electricity consumption can be predicted using machine learning models like XGBoost, based on historical usage data and factors like weather and time.