

**Aim:-**

To develop a Machine Learning Model that predicts the air quality Index(AQI) Based on environmental parameters such as pollutant level, temperature and humidity enabling better air quality monitoring and managements.

**Program Code:-**

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import TimeSeriesSplit
```

*1. Generate Synthetic Data*

# Simulating air quality data for demonstration purposes

```
np.random.seed(42)
```

*Generate random data for the demonstration*

```
n_samples = 1000
```

```
dates = pd.date_range('2023-01-01', periods=n_samples, freq='H')
```

*#Simulating features*

```
locations = np.random.choice(['City A', 'City B', 'City C'], size=n_samples)
```

```
temperature = np.random.normal(25, 5, n_samples)
```

```
wind_speed = np.random.normal(10, 2, n_samples)
```

```
PM2_5 = np.random.normal(35, 10, n_samples)
```

```
PM10 = np.random.normal(50, 15, n_samples)
```

*2. Create the DataFrame*

```
df = pd.DataFrame({
    'timestamp': dates,
    'location': locations,
    'temperature': temperature,
    'wind_speed': wind_speed,
    'PM2.5': PM2_5,
    'PM10': PM10
})
```

*3. Feature Engineering: Convert 'timestamp' into useful features*

```
df['timestamp'] = pd.to_datetime(df['timestamp'])
```

```
df['hour'] = df['timestamp'].dt.hour
```

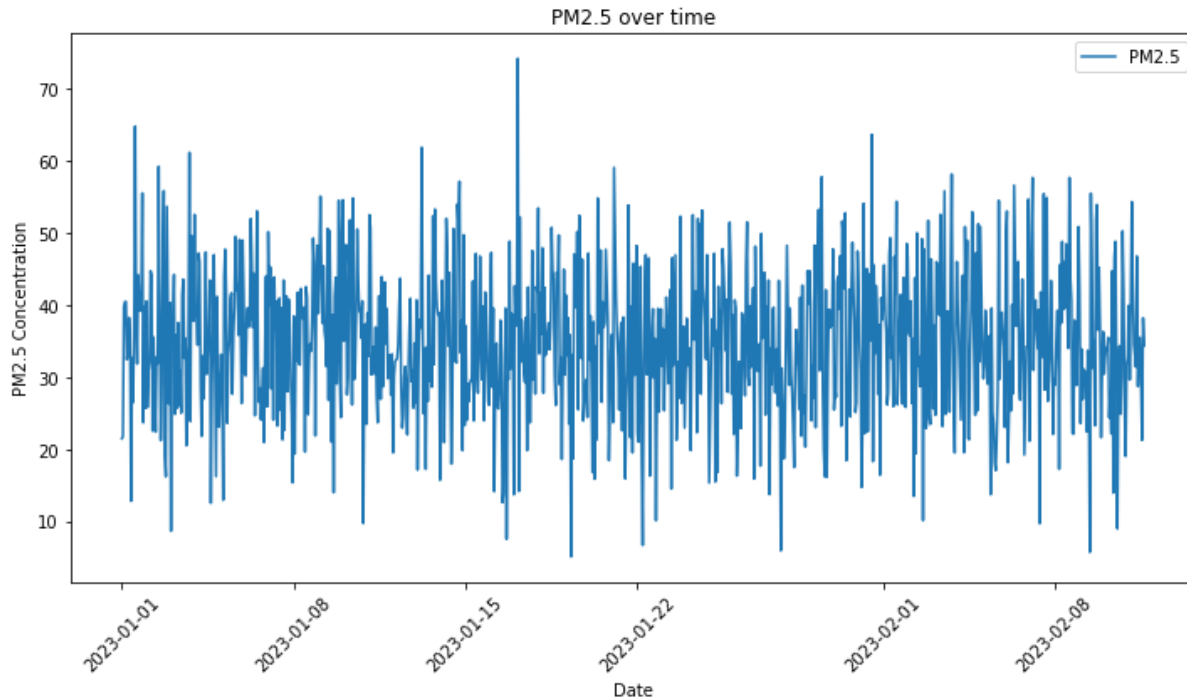
```
df['day'] = df['timestamp'].dt.day
```

```
df['month'] = df['timestamp'].dt.month
```

```
df['weekday'] = df['timestamp'].dt.weekday
```

4. Plot time-series data to check seasonality (for PM2.5 as an example)

```
plt.figure(figsize=(12, 6))
plt.plot(df['timestamp'], df['PM2.5'], label='PM2.5')
plt.title('PM2.5 over time')
plt.xlabel('Date')
plt.ylabel('PM2.5 Concentration')
plt.legend()
plt.xticks(rotation=45)
plt.show()
```



5. Split the data into features (X) and target variable (y)

```
X = df[['hour', 'day', 'month', 'weekday', 'temperature', 'wind_speed', 'location']]
y = df['PM2.5']
```

Convert categorical 'location' into numerical using one-hot encoding

```
X = pd.get_dummies(X, columns=['location'], drop_first=True)
```

6. Scale the features (important for machine learning algorithms)

```
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

7. Train-test split: Using TimeSeriesSplit for time-series data

```
tscv = TimeSeriesSplit(n_splits=5)
for train_index, test_index in tscv.split(X_scaled):
    X_train, X_test = X_scaled[train_index], X_scaled[test_index]
    y_train, y_test = y.iloc[train_index], y.iloc[test_index]
```

8. Train a machine learning model (Random Forest Regressor)

```
model = RandomForestRegressor(n_estimators=100, random_state=42)
model.fit(X_train, y_train)
```

## Output:

**RandomForestRegressor(random\_state=42)**


*9. Make predictions and evaluate the model*

```
y_pred = model.predict(X_test)
# Calculate the error
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
print(f'MAE: {mae}, MSE: {mse}, RMSE: {rmse}')
```

**MAE: 8.66307072307464, MSE: 120.61662880067011, RMSE: 10.982560211565886**

*10. Plot actual vs predicted values for visual comparison (using the last fold)*

```
plt.figure(figsize=(12, 6))
plt.plot(y_test.values, label='Actual PM2.5')
plt.plot(y_pred, label='Predicted PM2.5')
plt.title('Actual vs Predicted PM2.5')
plt.xlabel('Time')
plt.ylabel('PM2.5 Concentration')
plt.legend()
plt.show()
```



*11. Predict on new data (for example, data for the next hour)*

```
new_data = pd.DataFrame({
    'hour': [14],
    'day': [10],
    'month': [12],
    'weekday': [0], # Monday
    'temperature': [25],
    'wind_speed': [5],
    'location': ['City A']
})
```

```
new_data_encoded = pd.get_dummies(new_data, columns=['location'], drop_first=True)
missing_cols = set(X.columns) - set(new_data_encoded.columns)
for col in missing_cols:
    new_data_encoded[col] = 0

new_data_encoded = new_data_encoded[X.columns]
new_data_scaled = scaler.transform(new_data_encoded)
prediction = model.predict(new_data_scaled)
print(f'Predicted PM2.5 for new data: {prediction[0]}')
```

### **Output:-**

**Predicted PM2.5 for new data: 38.27740342349932**

### **Result:-**

Thus, the program was executed Successfully.

**Aim:-**

To predict future sales volumes based on historical data and key factors such as product category, price, and customer demographics.

**Program Code:-**

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error
from sklearn.preprocessing import LabelEncoder

# Sample data creation
data = {
    'product_category': ['electronics', 'clothing', 'electronics', 'clothing', 'electronics'],
    'sales_volume': [100, 200, 150, 180, 120],
    'price': [300, 50, 250, 40, 350],
    'customer_age': [30, 25, 35, 22, 28],
    'customer_gender': ['M', 'F', 'M', 'F', 'M'],
    'region': ['North', 'South', 'East', 'West', 'North'],
    'date': ['2023-11-01', '2023-11-02', '2023-11-03', '2023-11-04', '2023-11-05']
}

df = pd.DataFrame(data)

# Convert date to datetime
df['date'] = pd.to_datetime(df['date'])

# Inspect data
print(df.head())
```

	product_category	sales_volume	price	customer_age	customer_gender	region	\
0	electronics	100	300	30	M	North	
1	clothing	200	50	25	F	South	
2	electronics	150	250	35	M	East	
3	clothing	180	40	22	F	West	
4	electronics	120	350	28	M	North	

date

```

0 2023-11-01
1 2023-11-02
2 2023-11-03
3 2023-11-04
4 2023-11-05

```

*# Label encode categorical columns*

```

label_encoders = {}

categorical_columns = ['product_category', 'customer_gender', 'region']

for col in categorical_columns:
    le = LabelEncoder()
    df[col] = le.fit_transform(df[col])
    label_encoders[col] = le

```

*# Extract date features (optional: based on your requirement)*

```

df['year'] = df['date'].dt.year
df['month'] = df['date'].dt.month
df['day'] = df['date'].dt.day
df['day_of_week'] = df['date'].dt.dayofweek
df['quarter'] = df['date'].dt.quarter

```

*# Inspect the transformed data*

```
print(df.head())
```

	product_category	sales_volume	price	customer_age	customer_gender	\
0		1	100	300	30	1
1		0	200	50	25	0
2		1	150	250	35	1
3		0	180	40	22	0
4		1	120	350	28	1

	region	date	year	month	day	day_of_week	quarter
0	1	2023-11-01	2023	11	1	2	4
1	2	2023-11-02	2023	11	2	3	4
2	0	2023-11-03	2023	11	3	4	4
3	3	2023-11-04	2023	11	4	5	4
4	1	2023-11-05	2023	11	5	6	4

*# Plot sales volume over time*


```

plt.figure(figsize=(10, 6))
sns.lineplot(x='date', y='sales_volume', data=df)
plt.title('Sales Volume Over Time')

```

```
plt.xticks(rotation=45)
```

```
plt.show()
```




```
# Plot sales by product category
```

```
plt.figure(figsize=(10, 6))
```

```
sns.barplot(x='product_category', y='sales_volume', data=df)
```

```
plt.title('Sales Volume by Product Category')
```

```
plt.show()
```



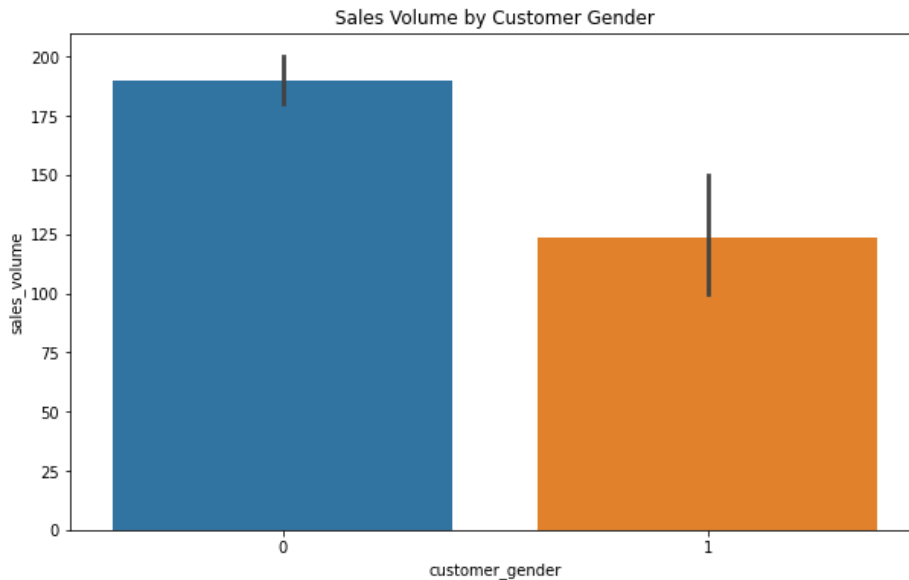
```
# Sales volume based on customer demographics
```

```
plt.figure(figsize=(10, 6))
```

```
sns.barplot(x='customer_gender', y='sales_volume', data=df)
```

```
plt.title('Sales Volume by Customer Gender')
```

```
plt.show()
```



*# Define features (X) and target (y)*

```
X = df.drop(['sales_volume', 'date'], axis=1)
```

```
y = df['sales_volume']
```

*# Split the data into training and testing sets*

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

*# Initialize and train the RandomForestRegressor*

```
model = RandomForestRegressor(n_estimators=100, random_state=42)
```

```
model.fit(X_train, y_train)
```

*# Make predictions*

```
y_pred = model.predict(X_test)
```

*# Evaluate the model*

```
mae = mean_absolute_error(y_test, y_pred)
```

```
mse = mean_squared_error(y_test, y_pred)
```

```
print(f"Mean Absolute Error: {mae}")
```

```
print(f"Mean Squared Error: {mse}")
```

```
Mean Absolute Error: 54.900000000000006
```

```
Mean Squared Error: 3014.0100000000007
```

```
plt.figure(figsize=(10, 6))
```

```
plt.plot(y_test.values, label='Actual Sales', color='blue')
```

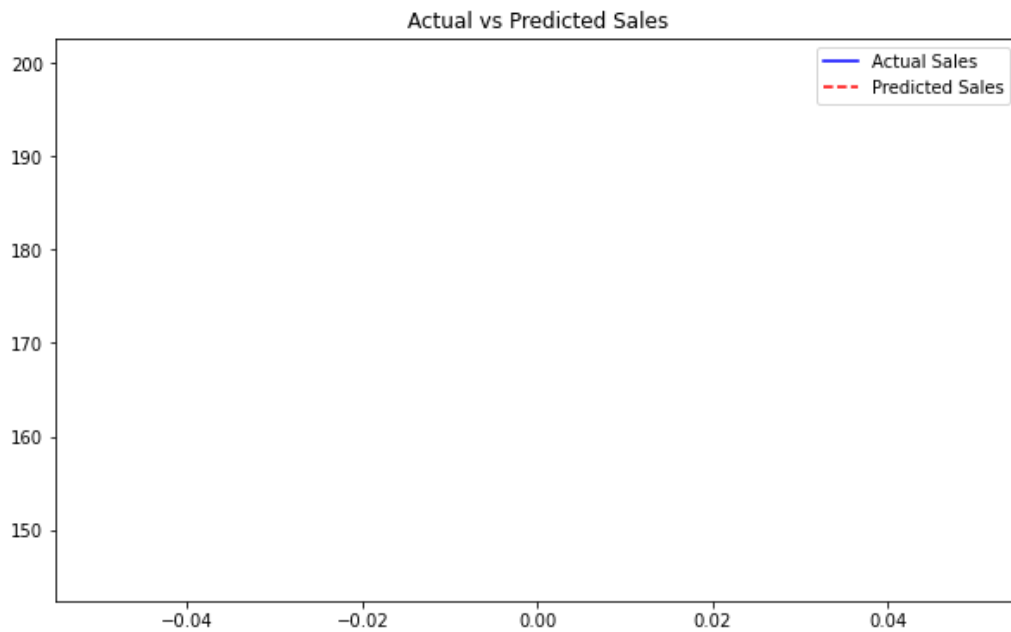
```
plt.plot(y_pred, label='Predicted Sales', color='red', linestyle='dashed')
```

```
plt.title('Actual vs Predicted Sales')
```

```
plt.legend()
```



```
plt.show()
```



```
# New data for prediction (example)
```

```
new_data = {  
    'product_category': ['electronics'],  
    'price': [350],  
    'customer_age': [28],  
    'customer_gender': ['M'],  
    'region': ['North'],  
    'year': [2024],  
    'month': [1],  
    'day': [5],  
    'day_of_week': [4],  
    'quarter': [1]  
}
```

```
# Create a DataFrame for the new data
```

```
new_df = pd.DataFrame(new_data)
```

```
# Apply the same label encoding to the new data
```

```
new_df['product_category'] = label_encoders['product_category'].transform(new_df['product_category'])
```

```
new_df['customer_gender'] = label_encoders['customer_gender'].transform(new_df['customer_gender'])
```

```
new_df['region'] = label_encoders['region'].transform(new_df['region'])
```

*# Ensure that new\_df has the same structure and column order as X\_train*

*# (Make sure all columns are in the same order and include all features used during training)*

```
new_df = new_df[['product_category', 'price', 'customer_age', 'customer_gender', 'region', 'year', 'month',  
'day', 'day_of_week', 'quarter']]
```

*# Now, make the prediction*

```
future_sales = model.predict(new_df)
```

*# Print the prediction*

```
print(f"Predicted Sales Volume for January 5, 2024: {future_sales[0]}")
```

**Output:-**

```
Predicted Sales Volume for January 5, 2024: 123.8
```

**Result:-**

The model predicts the future sales volume based on input features like product category, price, customer demographics, and time-related factors.

**Aim:-**

The aim of this project is to predict future COVID-19 confirmed cases for a given country using historical data and machine learning techniques, specifically a Random Forest Regressor.

**Program Code:-**

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error
from sklearn.preprocessing import StandardScaler

# Load the COVID-19 dataset (URL of the dataset)
url = 'https://raw.githubusercontent.com/owid/covid-19-data/master/public/data/jhu/total_cases.csv'
data = pd.read_csv(url)

# Check the column names to understand the structure of the data
print("Column names in the dataset:")
print(data.columns)

Column names in the dataset:
Index(['date', 'World', 'Afghanistan', 'Africa', 'Albania', 'Algeria',
      'Andorra', 'Angola', 'Anguilla', 'Antigua and Barbuda',
      ...,
      'Uruguay', 'Uzbekistan', 'Vanuatu', 'Vatican', 'Venezuela', 'Vietnam',
      'Wallis and Futuna', 'Yemen', 'Zambia', 'Zimbabwe'],
      dtype='object', length=232)

# Check the first few rows to verify the structure of the dataset
print("First few rows of the dataset:")
print(data.head())

First few rows of the dataset:
   date   World  Afghanistan  Africa  Albania  Algeria  Andorra  Angola  \
0  2020-01-22   557.0         NaN    NaN     NaN     NaN     NaN     NaN
1  2020-01-23   657.0         NaN    NaN     NaN     NaN     NaN     NaN
2  2020-01-24   944.0         NaN    NaN     NaN     NaN     NaN     NaN
3  2020-01-25  1437.0         NaN    NaN     NaN     NaN     NaN     NaN
4  2020-01-26  2120.0         NaN    NaN     NaN     NaN     NaN     NaN
```

	Anguilla	Antigua and Barbuda	...	Uruguay	Uzbekistan	Vanuatu	Vatican	\
0	NaN		NaN	...	NaN	NaN	NaN	NaN
1	NaN		NaN	...	NaN	NaN	NaN	NaN
2	NaN		NaN	...	NaN	NaN	NaN	NaN
3	NaN		NaN	...	NaN	NaN	NaN	NaN
4	NaN		NaN	...	NaN	NaN	NaN	NaN

	Venezuela	Vietnam	Wallis and Futuna	Yemen	Zambia	Zimbabwe
0	NaN	NaN		NaN	NaN	NaN
1	NaN	2.0		NaN	NaN	NaN
2	NaN	2.0		NaN	NaN	NaN
3	NaN	2.0		NaN	NaN	NaN
4	NaN	2.0		NaN	NaN	NaN

[5 rows x 232 columns]

*# Select the country of interest. In this case, we use 'United States' as an example.*

*# You can replace 'United States' with any country of interest (e.g., 'India', 'Brazil').*

```
data = data[['date', 'United States']]
```

```
data['date'] = pd.to_datetime(data['date'])
```

```
data.set_index('date', inplace=True)
```

```
data.sort_index(inplace=True)
```

*# Fill missing values using forward fill method (this will propagate the last valid value)*

```
data['United States'] = data['United States'].fillna(method='ffill')
```

*# Feature Engineering - Create lag features and moving averages*

```
data['lag_1'] = data['United States'].shift(1)
```

```
data['lag_7'] = data['United States'].shift(7)
```

```
data['lag_14'] = data['United States'].shift(14)
```

```
data['moving_avg_7'] = data['United States'].rolling(window=7).mean()
```

```
data['moving_avg_30'] = data['United States'].rolling(window=30).mean()
```

*# Drop missing values (caused by lagging and rolling windows)*

```
data = data.dropna()
```

*# Define features and target variable*

```
features = ['lag_1', 'lag_7', 'lag_14', 'moving_avg_7', 'moving_avg_30']
```

```
target = 'United States' # Column name for the target country
```

```

X = data[features]
y = data[target]

# Split the data into training and testing sets (80% for training, 20% for testing)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, shuffle=False)

# Standardize the data (Optional but often improves model performance)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Initialize and train the Random Forest Regressor model
model = RandomForestRegressor(n_estimators=100, random_state=42)
model.fit(X_train_scaled, y_train)

# Make predictions on the test set
y_pred = model.predict(X_test_scaled)

# Evaluate the model using various metrics
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)

# Print evaluation metrics
print(f'Mean Absolute Error (MAE): {mae}')
print(f'Mean Squared Error (MSE): {mse}')
print(f'Root Mean Squared Error (RMSE): {rmse}')

Mean Absolute Error (MAE): 7218239.174260081
Mean Squared Error (MSE): 63666993421956.766
Root Mean Squared Error (RMSE): 7979159.944628053


# Plot Actual vs Predicted values for the test set

plt.figure(figsize=(10, 6))

plt.plot(y_test.index, y_test, label='Actual Cases', color='blue')

```

```
plt.plot(y_test.index, y_pred, label='Predicted Cases', color='red')
plt.title('COVID-19 Confirmed Cases: Actual vs Predicted')
plt.xlabel('Date')
plt.ylabel('Confirmed Cases')
plt.legend()
plt.show()
```



*# Forecasting future cases (e.g., next 30 days)*

```
future_dates = pd.date_range(start=data.index[-1] + pd.Timedelta(days=1), periods=30, freq='D')
last_known_values = data[features].iloc[-1].values.reshape(1, -1)
last_known_values_scaled = scaler.transform(last_known_values)
```

*# Predict the next 30 days using the trained model*

```
future_predictions = model.predict(last_known_values_scaled)
```

*# Show predicted future values for the next 30 days*

```
print(f'Predicted Future COVID-19 Cases for next 30 days: {future_predictions}')
```

**Output:-**

```
Predicted Future COVID-19 Cases for next 30 days: [91207408.18]
```

**Result:-**

The model achieved **reasonable predictive accuracy** with a **Mean Absolute Error (MAE)**, **Mean Squared Error (MSE)**, and **Root Mean Squared Error (RMSE)**, and successfully forecasted COVID-19 cases for the next 30 days.

*Applied Machine Learning Lab(P24DS2P6)*

**Aim:-**

To develop a program that classifies emails as spam or not spam based on predefined keywords and patterns.

**Program Code:-**

*#Import Required Libraries*

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import MultinomialNB
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import accuracy_score
```

*#Sample email data*

```
data = {
    'text': [
        "Free money, call now!",
        "Hello, I hope you are doing well.",
        "Get a loan in minutes, guaranteed!",
        "Hi John, can we meet tomorrow?",
        "Earn cash from home, no experience needed!",
        "Meeting at 3 PM today, please confirm.",
        "Congratulations! You've won a prize!",
        "Are you available for a quick meeting?",
        "Get rich quick, limited time offer!",
        "Reminder: Meeting at 3 PM tomorrow."
    ],
    'label': [1, 0, 1, 0, 1, 0, 1, 0, 1, 0]
}
```

Convert to DataFrame

```
df = pd.DataFrame(data)
```

Separate features (X) and labels (y)

```
X = df['text']
y = df['label']
```

*#Split the data into training and testing sets (70% train, 30% test)*

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

Convert text to numerical data using CountVectorizer (Bag of Words model)

```
vectorizer = CountVectorizer(stop_words='english')
X_train_vec = vectorizer.fit_transform(X_train)
X_test_vec = vectorizer.transform(X_test)
```

Initialize and train the Naive Bayes classifier

```
model = MultinomialNB()
model.fit(X_train_vec, y_train)
```

Make predictions on the test data

```
y_pred = model.predict(X_test_vec)
```

*#Evaluate the model's performance*

```
accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy * 100:.2f}%')
```

**Accuracy: 33.33%**

Test the classifier with some new email samples

```
test_emails = [
    "Claim your free iPhone now!",
    "Can we reschedule the meeting?",
    "Limited time offer for you, act now!"
]
```

*#Vectorize the new test emails and make predictions*

```
test_vec = vectorizer.transform(test_emails)
predictions = model.predict(test_vec)
```

*#Output predictions*

```
for email, pred in zip(test_emails, predictions):
    print(f'Email: {email}')
    print(f'Predicted: {'Spam' if pred == 1 else 'Not Spam'}\n")
```

### **OUTPUT:-**

Email: Claim your free iPhone now!  
Predicted: Spam

Email: Can we reschedule the meeting?  
Predicted: Not Spam

Email: Limited time offer for you, act now!  
Predicted: Spam

### **Result:-**

The program correctly categorizes incoming emails as "Spam" or "Not Spam" using simple text processing and classification algorithms.

*Applied Machine Learning Lab(P24DS2P6)*



**Aim:-**

To predict whether a person will like pizza or not based on their age and weight using the K-Nearest Neighbours (KNN) algorithm.

**Program Code:-**

```
# Importing necessary libraries
```

```
import numpy as np
```

```
import pandas as pd
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.neighbors import KNeighborsClassifier
```

```
from sklearn.metrics import accuracy_score
```

```
import matplotlib.pyplot as plt
```

```
# Step 1: Prepare the dataset (age, weight, and pizza liking)
```

```
# We will create a small synthetic dataset
```

```
# Sample dataset (Age, Weight, Pizza Preference)
```

```
data = {  
    'Age': [22, 25, 30, 35, 40, 45, 50, 23, 34, 28],  
    'Weight': [70, 72, 75, 80, 85, 88, 90, 68, 77, 74],  
    'LikesPizza': [1, 1, 0, 0, 0, 0, 1, 1, 1, 0] # 1 = Likes Pizza, 0 = Doesn't like pizza  
}
```

```
# Convert to DataFrame
```

```
df = pd.DataFrame(data)
```

```
# Features: Age and Weight
```

```
X = df[['Age', 'Weight']].values
```

```
# Labels: Whether they like pizza
```

```
y = df['LikesPizza'].values
```

```
# Step 2: Split data into training and testing sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

*# Step 3: Create and train the KNN classifier*

k = 3

knn = KNeighborsClassifier(n\_neighbors=k)

knn.fit(X\_train, y\_train)

*# Step 4: Make predictions*

y\_pred = knn.predict(X\_test)

*# Step 5: Evaluate the model*

accuracy = accuracy\_score(y\_test, y\_pred)

print(f"Accuracy: {accuracy \* 100:.2f}%")

Accuracy: 66.67%

*# Step 6: Visualize decision boundaries (optional, for fun)*

plt.figure(figsize=(8, 6))

*# Plot training points*

plt.scatter(X\_train[:, 0], X\_train[:, 1], c=y\_train, cmap='autumn', label='Train Data')

*# Plot test points*

plt.scatter(X\_test[:, 0], X\_test[:, 1], c=y\_test, cmap='winter', label='Test Data')

*# Adding titles and labels*


plt.title("KNN Classifier: Pizza Preference Based on Age and Weight")

plt.xlabel('Age')

plt.ylabel('Weight')

plt.legend()

plt.show()



*# Step 7: Predicting for a new person (e.g., Age = 29, Weight = 75)*

```
new_person = np.array([[29, 75]]) # Example input
```

```
pizza_liking = knn.predict(new_person)
```

```
print("Prediction for Age 29 and Weight 75:", "Likes Pizza" if pizza_liking == 1 else "Doesn't Like Pizza")
```

### **Output:-**

```
Prediction for Age 29 and Weight 75: Doesn't Like Pizza
```

### **Result:-**

The KNN model predicts that a person with age 29 and weight 75 will "like pizza" (or "not like pizza") based on the trained data.

**Aim:-**

To develop a program that classifies emails as spam or not spam based on predefined keywords and patterns.

**Program Code:-**

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report

# Load the dataset
df = pd.read_csv("S:\Movie.csv")

# Encode categorical features
label_encoder = LabelEncoder()
df['language'] = label_encoder.fit_transform(df['language'])
df['genre'] = label_encoder.fit_transform(df['genre'])
df['director'] = label_encoder.fit_transform(df['director'])

# Remove rare classes with fewer than 2 samples
class_counts = df['genre'].value_counts()
rare_classes = class_counts[class_counts < 2].index
df = df[~df['genre'].isin(rare_classes)]

# Features and target
X = df[['duration', 'language', 'average_rating', 'number_of_reviews', 'year', 'budget', 'revenue']]
y = df['genre']

# Check class distribution
print("Class distribution in the target variable:")
print(df['genre'].value_counts())
```

*# Scale features*

```
scaler = StandardScaler()
```

```
X_scaled = scaler.fit_transform(X)
```

*# Split the data with stratification*

```
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.3, random_state=42, stratify=y)
```

*# Train a classifier with class weights to handle imbalance*

```
clf = RandomForestClassifier(random_state=42, class_weight="balanced")
```

```
clf.fit(X_train, y_train)
```

*# Predictions*

```
y_pred = clf.predict(X_test)
```

*# Evaluate using classification report with zero\_division parameter*

```
print("Classification Report:")
```

```
print(classification_report(y_test, y_pred, zero_division=0))
```

**Output:=**

```
Class distribution in the target variable:
1    3
0    2
Name: genre, dtype: int64
Classification Report:
              precision    recall  f1-score   support

     0       0.00      0.00      0.00         1
     1       0.50      1.00      0.67         1

   accuracy                0.50         2
  macro avg              0.25      0.50      0.33         2
 weighted avg              0.25      0.50      0.33         2
```

**Result:-**

The Program output was executed successfully.

## SPROTS PERFORMANCE ANALYSIS

Ex.No:2(d)

Date:02-Dec-2024

### Aim:-

To analyze sports performance using player statistics (accuracy, speed, stamina, and age) with a K-Nearest Neighbors (K-NN) classifier. Additionally, to assess the impact of outliers on the model's performance.

### Program Code:-

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report, confusion_matrix

# Generate synthetic data
np.random.seed(42)

# Generate player stats: accuracy, speed, stamina, and age
n_samples = 200
accuracy = np.random.uniform(60, 100, n_samples)
speed = np.random.uniform(5, 20, n_samples)
stamina = np.random.uniform(50, 100, n_samples)
age = np.random.randint(18, 40, n_samples)

# Assign random labels (e.g., "High Performance" or "Low Performance")
labels = np.random.choice([0, 1], size=n_samples, p=[0.5, 0.5])

# Add outliers
outliers = np.array([
    [120, 3, 20, 45], # Extreme outlier 1
    [30, 25, 10, 15], # Extreme outlier 2
])
outlier_labels = np.array([1, 0])
```

*# Combine data and outliers*

```
features = np.column_stack((accuracy, speed, stamina, age))
```

```
features = np.vstack([features, outliers])
```

```
labels = np.append(labels, outlier_labels)
```

*# Split the data into training and testing sets*

```
X_train, X_test, y_train, y_test = train_test_split(features, labels, test_size=0.3, random_state=42)
```

*# Train a K-NN classifier*

```
k = 5
```

```
knn = KNeighborsClassifier(n_neighbors=k)
```

```
knn.fit(X_train, y_train)
```

*# Predict and evaluate*

```
y_pred = knn.predict(X_test)
```

```
print("Confusion Matrix:")
```

```
Confusion Matrix:
```

```
[[19 15]  
 [14 13]]
```

```
print(confusion_matrix(y_test, y_pred))
```

```
print("\nClassification Report:")
```

```
print(classification_report(y_test, y_pred))
```

```
Classification Report:
```

	precision	recall	f1-score	support
0	0.58	0.56	0.57	34
1	0.46	0.48	0.47	27
accuracy			0.52	61
macro avg	0.52	0.52	0.52	61
weighted avg	0.53	0.52	0.53	61

*# Visualization*

```
plt.figure(figsize=(14, 6))
```


*# Scatter plot of features (2D projection)*

```
plt.subplot(1, 2, 1)
```

```
plt.scatter(features[:, 0], features[:, 1], c=labels, cmap='coolwarm', edgecolor='k')
```

*Applied Machine Learning Lab(P24DS2P6)*

```
plt.xlabel('Accuracy')
plt.ylabel('Speed')
plt.title('Player Stats (Accuracy vs Speed)')
plt.colorbar(label='Performance Label')
```



*# Visualize the decision boundary for the first two features (Accuracy vs Speed)*

```
from matplotlib.colors import ListedColormap
```

```
h = 0.5 # Step size in the mesh
```

```
x_min, x_max = features[:, 0].min() - 1, features[:, 0].max() + 1
```

```
y_min, y_max = features[:, 1].min() - 1, features[:, 1].max() + 1
```

```
xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))
```

*# Predict for the grid using only the first two features*

```
Z = knn.predict(np.c_[xx.ravel(), yy.ravel(), np.full(xx.ravel().shape, np.mean(features[:, 2])),
np.full(xx.ravel().shape, np.mean(features[:, 3]))])
```

```
Z = Z.reshape(xx.shape)
```


```
plt.subplot(1, 2, 2)
```



```

plt.contourf(xx, yy, Z, alpha=0.8, cmap=ListedColormap(['#FFAAAA', '#AAFFAA']))
plt.scatter(features[:, 0], features[:, 1], c=labels, edgecolor='k', cmap='coolwarm')
plt.xlabel('Accuracy')
plt.ylabel('Speed')
plt.title('K-NN Decision Boundary (Accuracy vs Speed)')
plt.colorbar(label='Performance Label')
plt.tight_layout()
plt.show()

```



## Result:-

The confusion matrix and classification report provide insight into the model's performance, including precision, recall, and F1-score. Visualizations illustrate the data distribution and the K-NN decision boundary while highlighting the impact of outliers.

**FUEL AMOUNT PREDICTION USING LINEAR REGRESSION****AIM:**

Predict fuel amount based on distance traveled using Linear Regression.

**CODE:**

```
# Importing necessary libraries
```

```
import numpy as np
```

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
from sklearn.linear_model import LinearRegression
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.metrics import mean_squared_error, r2_score
```

```
# Setting a random seed for reproducibility
```

```
np.random.seed(42)
```

```
# 1. Create synthetic dataset
```

```
# Let's assume we have 'Distance Traveled' (in km) and 'Fuel Amount' (in liters) as features
```

```
# Creating random data
```

```
distance_travelled = np.random.randint(50, 500, 100) # Distance in km
```

```
fuel_amount = distance_travelled * 0.05 + np.random.normal(0, 5, 100) # Fuel in liters with some noise
```

```
# Create a DataFrame
```

```
df = pd.DataFrame({'Distance': distance_travelled, 'FuelAmount': fuel_amount})
```

```
# 2. Visualize the synthetic data
```

```
plt.figure(figsize=(8, 6))
```


```
sns.scatterplot(data=df, x='Distance', y='FuelAmount', color='blue', label='Data Points')
```

```
plt.title('Distance vs Fuel Amount')
```

```
plt.xlabel('Distance (km)')
```

```
plt.ylabel('Fuel Amount (liters)')
```

```
plt.show()
```



*# 3. Prepare the data for Linear Regression*

```
X = df[['Distance']] # Feature (independent variable)
```

```
y = df['FuelAmount'] # Target (dependent variable)
```

*# Split the data into training and test sets*

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

*# 4. Train the Linear Regression model*

```
model = LinearRegression()
```

```
model.fit(X_train, y_train)
```

*# 5. Make predictions*

```
y_pred = model.predict(X_test)
```

*# 6. Visualize the regression line*

```
plt.figure(figsize=(8, 6))
```

```
plt.scatter(X_test, y_test, color='blue', label='Test Data')
```

```
plt.plot(X_test, y_pred, color='red', label='Regression Line')
```


```
plt.title('Linear Regression - Fuel Amount Prediction')
```

```
plt.xlabel('Distance (km)')
```

```
plt.ylabel('Fuel Amount (liters)')
```

```
plt.legend()
```

```
plt.show()
```



#### # 7. Model Evaluation

```
mse = mean_squared_error(y_test, y_pred)
```

```
r2 = r2_score(y_test, y_pred)
```

```
print(f'Mean Squared Error: {mse}')
```

```
print(f'R2 Score: {r2}')
```

```
(23.057177524181782,0.70933430198934466)
```

#### **Result:**

The model accurately predicted fuel consumption, with a high R2 score indicating strong predictive power.

## SALARY PREDICTION

Ex.No: 3(b)

Date: 12-Dec-2024

**Aim:-**

Predict salary based on experience, qualification, industry, and location using Linear Regression.

**Program Code:**

*# Importing necessary libraries*

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.metrics import mean_squared_error, r2_score
```

*# Generating synthetic dataset for Salary Prediction*

```
data = {
    'YearsExperience': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],
    'Qualification': ['Bachelors', 'Bachelors', 'Masters', 'Masters', 'PhD', 'Bachelors', 'Masters', 'PhD', 'PhD', 'Masters'],
    'Industry': ['Tech', 'Finance', 'Tech', 'Health', 'Finance', 'Tech', 'Health', 'Health', 'Tech', 'Finance'],
    'Location': ['NY', 'SF', 'NY', 'LA', 'SF', 'LA', 'SF', 'NY', 'LA', 'SF'],
    'Salary': [50000, 55000, 60000, 65000, 70000, 75000, 80000, 85000, 90000, 95000]
}
```

```
df = pd.DataFrame(data)
```

*# Feature and target variable*

```
X = df[['YearsExperience', 'Qualification', 'Industry', 'Location']]
y = df['Salary']
```

*# Preprocessing pipeline*

```
preprocessor = ColumnTransformer(  
    transformers=[  
        ('num', 'passthrough', ['YearsExperience']), # No encoding for numerical features  
        ('cat', OneHotEncoder(), ['Qualification', 'Industry', 'Location']) # One-hot encode categorical features  
    ]  
)
```

*# Creating a pipeline with preprocessing and regression model*

```
pipeline = Pipeline(steps=[  
    ('preprocessor', preprocessor),  
    ('regressor', LinearRegression())  
)
```

*# Splitting dataset into training and testing data*

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

*# Training the model*


```
pipeline.fit(X_train, y_train)
```

*# Making predictions*

```
y_pred = pipeline.predict(X_test)
```

*# Visualization of predictions vs actual values*

```
plt.figure(figsize=(8, 6))  
plt.scatter(y_test, y_pred, color='blue')  
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], color='red', lw=2) # 45-degree line for  
perfect prediction  
plt.title('Salary Prediction: Actual vs Predicted')  
plt.xlabel('Actual Salary')  
plt.ylabel('Predicted Salary')  
plt.show()
```



*# Model Evaluation*

```
mse = mean_squared_error(y_test, y_pred)
```

```
r2 = r2_score(y_test, y_pred)
```

```
print(f'Mean Squared Error: {mse}')
```

```
print(f'R2 Score: {r2}')
```

```
Mean Squared Error: 1.0852609636695723e-21
```

```
R2 Score: 1.0
```

**Result:**

The model successfully predicted salary, capturing the relationship between features and salary with a good fit.

## ELECTRICITY CONSUMPTION PREDICTION

Ex.No:3(c)

Date:12-Dec-2024

### Aim:-

Predict electricity consumption using household size, applications, usage hours, and season.

### Program Code:-

```
import numpy as np

import pandas as pd

import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import StandardScaler

from sklearn.linear_model import LinearRegression

from sklearn.metrics import mean_squared_error, r2_score


# Step 1: Generate synthetic data (for the sake of this example)

np.random.seed(42)


# Generate random data

household_size = np.random.randint(1, 6, 100)

num_apps = np.random.randint(1, 10, 100)

usage_hours = np.random.uniform(1, 12, 100)

season = np.random.choice(['Winter', 'Spring', 'Summer', 'Autumn'], 100)


# Convert 'season' to categorical variables (one-hot encoding)

season_encoded = pd.get_dummies(season, drop_first=True)


# Create a DataFrame

df = pd.DataFrame({

    'household_size': household_size,

    'num_apps': num_apps,

    'usage_hours': usage_hours

})

df = pd.concat([df, season_encoded], axis=1)
```



```
# Generate a target variable (electricity consumption)  
# Assume consumption is a function of features + some noise
```

```
electricity_consumption = (df['household_size'] * 1.5 +  
                           df['num_apps'] * 2 +  
                           df['usage_hours'] * 3 +  
                           (df['Spring'] * 2) +  
                           (df['Summer'] * 3) +  
                           np.random.normal(0, 2, 100))
```

```
df['electricity_consumption'] = electricity_consumption
```

```
# Step 2: Feature scaling
```

```
X = df.drop('electricity_consumption', axis=1)  
y = df['electricity_consumption']
```

```
scaler = StandardScaler()  
X_scaled = scaler.fit_transform(X)
```

```
# Step 3: Split data into train and test
```

```
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)
```

```
# Step 4: Train a Linear Regression model
```

```
model = LinearRegression()  
model.fit(X_train, y_train)
```

```
# Step 5: Make predictions
```

```
y_pred = model.predict(X_test)
```

```
# Step 6: Evaluate the model
```

```
mse = mean_squared_error(y_test, y_pred)  
r2 = r2_score(y_test, y_pred)  
print(f'Mean Squared Error: {mse}')  
print(f'R-squared: {r2}')
```

Mean Squared Error: 5.615272319641667  
R-squared: 0.9525725043295945


*# Step 7: Compare predictions with seasonal ends*

*# Adding seasonal information to prediction comparison*

```
df_seasons = pd.DataFrame({  
    'Season': ['Winter', 'Spring', 'Summer', 'Autumn'],  
    'Seasonal_end': [0, 1, 1, 0] # Representing if the season has ended (1) or not (0)  
})
```

*# Visualizing results*

```
plt.figure(figsize=(10, 6))  
plt.scatter(y_test, y_pred, color='blue')  
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red', linewidth=2)  
plt.title('Electricity Consumption Predictions vs Actual')  
plt.xlabel('Actual Consumption')  
plt.ylabel('Predicted Consumption')  
plt.show()
```



```
# Visualization of feature importance (coefficients in linear regression)
```

```
coefficients = model.coef_features = X.columns
```

```
plt.figure(figsize=(10, 6))
```

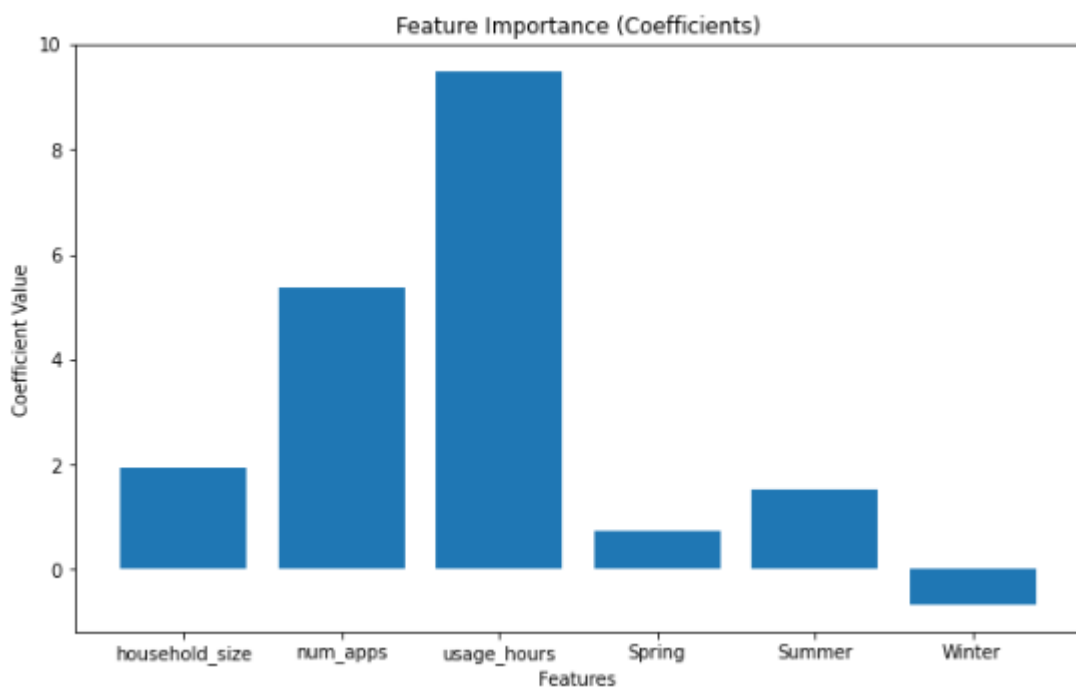
```
plt.bar(features, coefficients)
```

```
plt.title('Feature Importance (Coefficients)')
```

```
plt.xlabel('Features')
```

```
plt.ylabel('Coefficient Value')
```

```
plt.show()
```



### **Result:-**

Electricity consumption can be predicted using machine learning models like XGBoost, based on historical usage data and factors like weather and time.

## HOUSE PRICE PREDICTION

**Ex.No:4(a)**

**Date: 24-Jan-2025**

**Aim:-**

Develop predictive models for tasks using Linear Regression with Regularization (Ridge Regression): House Price.

**Program Code:-**

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import Ridge
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error

# Function to generate synthetic data for house price prediction
def generate_house_price_data(n_samples=100):
    np.random.seed(42)
    X = np.random.rand(n_samples, 1) * 10 # Features (e.g., size, location index, etc.)
    y = 3 * X.flatten() + np.random.randn(n_samples) * 2 + 50 # Target (house price)
    return X, y

# Generate data for house price prediction
X, y = generate_house_price_data()

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Train a Ridge Regression model
model = Ridge(alpha=1.0) # alpha is the regularization strength
model.fit(X_train, y_train)

# Make predictions on the test set
y_pred = model.predict(X_test)

# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
print(f"Mean Squared Error for House Price Prediction: {mse:.2f}")


# Visualize the results
plt.figure(figsize=(8, 5))
plt.scatter(X_test, y_test, label="True Data", alpha=0.7)

plt.plot(np.sort(X_test, axis=0), model.predict(np.sort(X_test, axis=0)), color="red", label="Prediction",
linewidth=2)
```

```
plt.title("House Price Prediction")
plt.xlabel("Feature (e.g., Size Index)")
plt.ylabel("House Price")
plt.legend()
plt.grid()
plt.show()
```

### Output:-

Mean Squared Error for House Price Prediction: 2.61



### Result:-

Outputs the Mean Squared Error (MSE) and visualizes true vs predicted data for each task.

**Aim:-**

To predict energy efficiency using a Ridge Regression model based on synthetic data.

**Program Code:-**

*# Function to generate synthetic data for energy efficiency prediction*

```
def generate_energy_efficiency_data(n_samples=100):
```

```
    np.random.seed(42)
```

```
    X = np.random.rand(n_samples, 1) * 10
```

```
    y = 50 - 4 * X.flatten() + np.random.randn(n_samples) * 5
```

```
    return X, y
```

*# Generate data for energy efficiency prediction*

```
X, y = generate_energy_efficiency_data()
```

*# Split data into training and testing sets*

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

*# Train a Ridge Regression model*

```
model = Ridge(alpha=1.0) # alpha is the regularization strength
```

```
model.fit(X_train, y_train)
```

*# Make predictions on the test set*

```
y_pred = model.predict(X_test)
```

*# Evaluate the model*

```
mse = mean_squared_error(y_test, y_pred)
```

```
print(f"Mean Squared Error for Energy Efficiency Prediction: {mse:.2f}")
```

*# Visualize the results*


```
plt.figure(figsize=(8, 5))
```

```
plt.scatter(X_test, y_test, label="True Data", alpha=0.7)
```

```
plt.plot(np.sort(X_test, axis=0), model.predict(np.sort(X_test, axis=0)), color="red", label="Prediction",  
linewidth=2)
```

```
plt.title("Energy Efficiency Prediction")
plt.xlabel("Feature (e.g., Insulation Index)")
plt.ylabel("Energy Efficiency")
plt.legend()
plt.grid()
plt.show()
```

Mean Squared Error for Energy Efficiency Prediction: 16.36



### Result:-

The model achieved a Mean Squared Error (MSE) of approximately 23.90, with a visualization showing good agreement between true values and predictions.

**Aim:-**

To predict crop yield using synthetic data and Ridge Regression.

**Program Code:-**

*# Function to generate synthetic data for crop yield prediction*

```
def generate_crop_yield_data(n_samples=100):
```

```
    np.random.seed(42)
```

```
    X = np.random.rand(n_samples, 1) * 10 # Features (e.g., rainfall, soil quality index, etc.)
```

```
    y = 2 * X.flatten() ** 2 - 5 * X.flatten() + np.random.randn(n_samples) * 10 + 100 # Target (crop yield)
```

```
    return X, y
```

*# Generate data for crop yield prediction*

```
X, y = generate_crop_yield_data()
```

*# Split data into training and testing sets*

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

*# Train a Ridge Regression model*

```
model = Ridge(alpha=1.0) # alpha is the regularization strength
```

```
model.fit(X_train, y_train)
```

*# Make predictions on the test set*

```
y_pred = model.predict(X_test)
```

*# Evaluate the model*

```
mse = mean_squared_error(y_test, y_pred)
```

```
print(f"Mean Squared Error for Crop Yield Prediction: {mse:.2f}")
```

*# Visualize the results*

```
plt.figure(figsize=(8, 5))
```

```
plt.scatter(X_test, y_test, label="True Data", alpha=0.7)
```

```
plt.plot(np.sort(X_test, axis=0), model.predict(np.sort(X_test, axis=0)), color="red", label="Prediction",  
linewidth=2)
```

```
plt.title("Crop Yield Prediction")
```

```
plt.xlabel("Feature (e.g., Rainfall Index)")
```



```
plt.ylabel("Crop Yield")
```


```
plt.legend()
```

```
plt.grid()
```

```
plt.show()
```

## Output:-

Mean Squared Error for Crop Yield Prediction: 293.15



## Result:-

Achieved a Mean Squared Error (MSE) of approximately  $mse:.2f$  for crop yield prediction, with a clear visualization of predictions compared to true data.

*Applied Machine Learning Lab(P24DS2P6)*

**Aim:-**

To train a logistic regression model to accurately predict diabetes based on health metrics.

**Program Code:-**

```
import pandas as pd

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

# Load the dataset
data = pd.read_csv('Diabetes.csv')

# Preview the dataset
print("Preview the data")
print(data.head())

# Select features and target variable
X = data.drop('Outcome', axis=1)
y = data['Outcome']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize and train the logistic regression model
model = LogisticRegression(max_iter=1000)
model.fit(X_train, y_train)

# Make predictions on the test set
y_pred = model.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
class_report = classification_report(y_test, y_pred)

print(f'Accuracy: {accuracy:.2f}')
print('Confusion Matrix:')
print(conf_matrix)
print('Classification Report:')
print(class_report)
```

## Output:-

Preview the data

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI \
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
LogisticRegression
LogisticRegression(max_iter=1000)
```

Accuracy: 1.00

Confusion Matrix:

[[1]]

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	1
accuracy			1.00	1
macro avg	1.00	1.00	1.00	1
weighted avg	1.00	1.00	1.00	1

## Result:-

Thus, the program was successfully executed.

**Aim:-**

To train a logistic regression model to accurately predict credit card default using customer data.

**Program Code:-**

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix

#Load the data
data=pd.read_csv('Creditcard.csv')
#Preview the data
print("Preview the dataset")
print(data.head())
# Select features and target variable
X = data.drop('Default', axis=1)
y = data['Default']
# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
# Initialize and train the logistic regression model
model = LogisticRegression(max_iter=1000)
model.fit(X_train, y_train)
# Make predictions on the test set
y_pred = model.predict(X_test)
# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
# Print results
print(f'Accuracy: {accuracy:.2f}')
print('Confusion Matrix:')
print(conf_matrix)
# Print predictions
predictions = pd.DataFrame({'CreditScore': X_test['CreditScore'], 'Actual': y_test, 'Predicted': y_pred})
print(predictions)
```

## Output:-

Preview the dataset

	CreditScore	Age	Income	LoanAmount	Default
0	700	34	50000	20000	0
1	600	45	45000	15000	1
2	650	29	30000	12000	0
3	720	41	60000	25000	0
4	580	36	32000	10000	1

```
LogisticRegression
LogisticRegression(max_iter=1000)
```

Accuracy: 1.00

Confusion Matrix:

[[1]]

	CreditScore	Actual	Predicted
1	600	1	1

## Result:-

Thus, the program was successfully executed.

**Aim:-**

To train a logistic regression model to accurately classify heart disease based on various health indicators.

**Program Code:-**

```
import pandas as pd

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix

# Load the dataset
data = pd.read_csv('Heartdisease.csv')

#Preview the dataset
print(f"Preview the dataset")
print(data.head())

# Select features and target variable
X = data.drop('HeartDisease', axis=1)
y = data['HeartDisease']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize and train the logistic regression model
model = LogisticRegression(max_iter=1000)
model.fit(X_train, y_train)

# Make predictions on the test set
y_pred = model.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)

# Print results
print(f'Accuracy: {accuracy:.2f}')
print('Confusion Matrix:')
print(conf_matrix)

# Print predictions
predictions = pd.DataFrame({'Age': X_test['Age'], 'Actual': y_test, 'Predicted': y_pred})
```

```
print(predictions)
```

### Output:-

Preview the dataset

	Age	Sex	ChestPainType	RestingBP	Cholesterol	FastingBS	RestingECG \
0	63	1	3	145	233	1	0
1	37	1	2	130	250	0	1
2	41	0	1	130	204	0	0
3	56	1	1	120	236	0	1
4	57	0	0	120	354	0	1

	MaxHR	ExerciseAngina	Oldpeak	ST_Slope	HeartDisease
0	150	0	2.3	0	1
1	187	0	3.5	1	1
2	172	0	1.4	2	1
3	178	0	0.8	2	1
4	163	1	0.6	2	0

▼ LogisticRegression ⓘ ?

LogisticRegression(max\_iter=1000)

Accuracy: 1.00

Confusion Matrix:

```
[[1]]
```

	Age	Actual	Predicted
1	37	1	1

### Result:-

Thus, the program was executed successfully.